

# NL2SQL (Text-to-SQL)

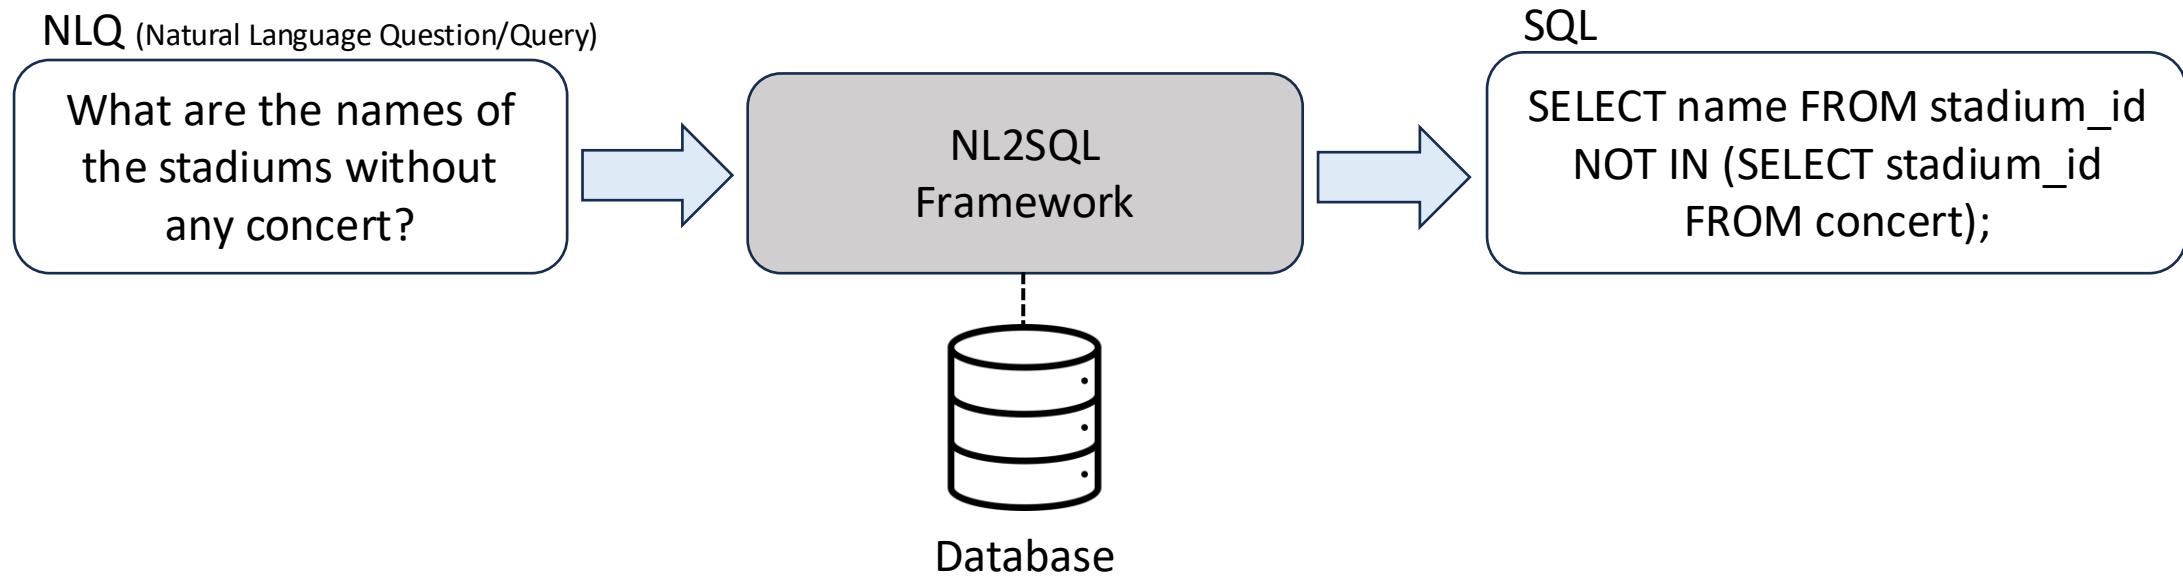
빅데이터 핀테크

Prof. Sang-Won Lee

TA. KyongShik Lee

## What is NL2SQL?

- **NL2SQL** (Natural Language to SQL) converts a user's natural-language question into an executable SQL query.
- It enables non-experts to query databases without writing SQL manually.
- The model must understand the question's intent and map it to the database schema (tables, columns, ..)



# Challenges of NL2SQL

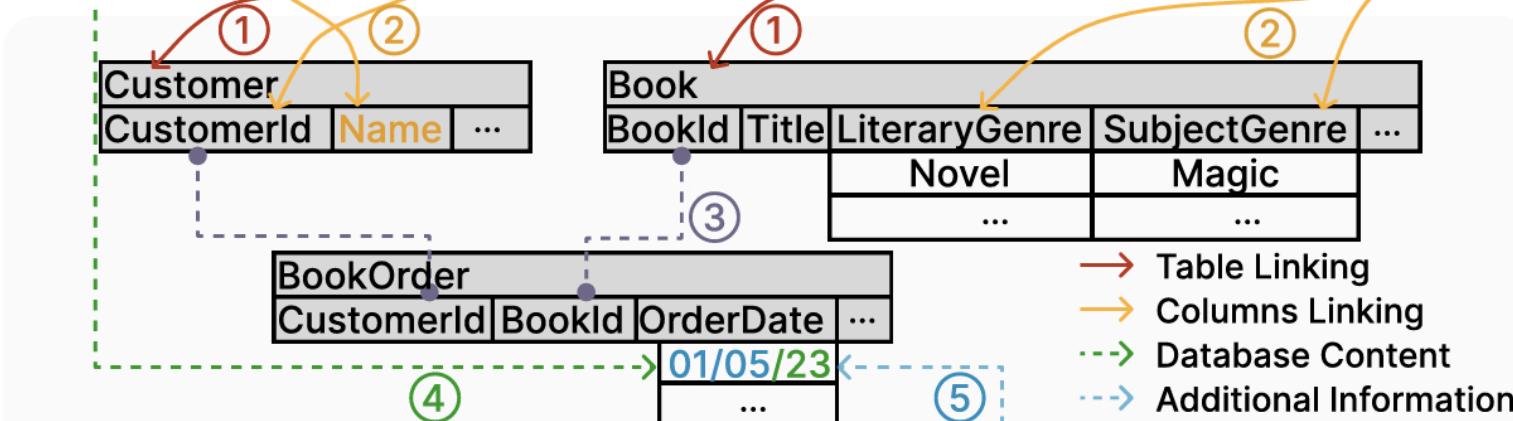
- **Uncertain natural-language queries**
  - User questions are often ambiguous or incomplete, making it difficult for the system to infer the exact intended meaning.
- **Complex databases and dirty content**
  - Real-world databases have complex relationships across tables, and column names or values can be inconsistent or contain errors, which makes query generation harder.
- **From natural language to SQL**
  - Translating flexible natural language into the constrained, formal structure of SQL is inherently challenging.
- **Engineering challenges in NL2SQL systems**
  - Practical systems must consider factors such as cost-effective deployment and model efficiency (latency, resource usage, scalability)

# Challenges of NL2SQL

## NL Query:

Find the **names** of all **customers** who checked out **books** on exactly 3 different **genres** on **Labor Day in 2023**.

## Database:



Additional Information: Note that **Labor Day stand for May 1**

## SQL:

```
SELECT Name
FROM Customer
NATURAL JOIN BookOrder
NATURAL JOIN Book
WHERE OrderDate='01/05/23'
GROUP BY CustomerId, Name
HAVING COUNT(DISTINCT SubjectGenre)=3
```

```
SELECT Name
FROM Customer
WHERE CustomerId = (SELECT CustomerId
FROM BookOrder NATURAL JOIN Book
WHERE OrderDate='01/05/23')
GROUP BY CustomerId
HAVING COUNT(DISTINCT SubjectGenre)=3
```

# Typical NL2SQL pipeline

## 1. Intent Understanding

- What does the user really want? (filters, aggregation, time range, top-k, etc)

## 2. Schema Linking

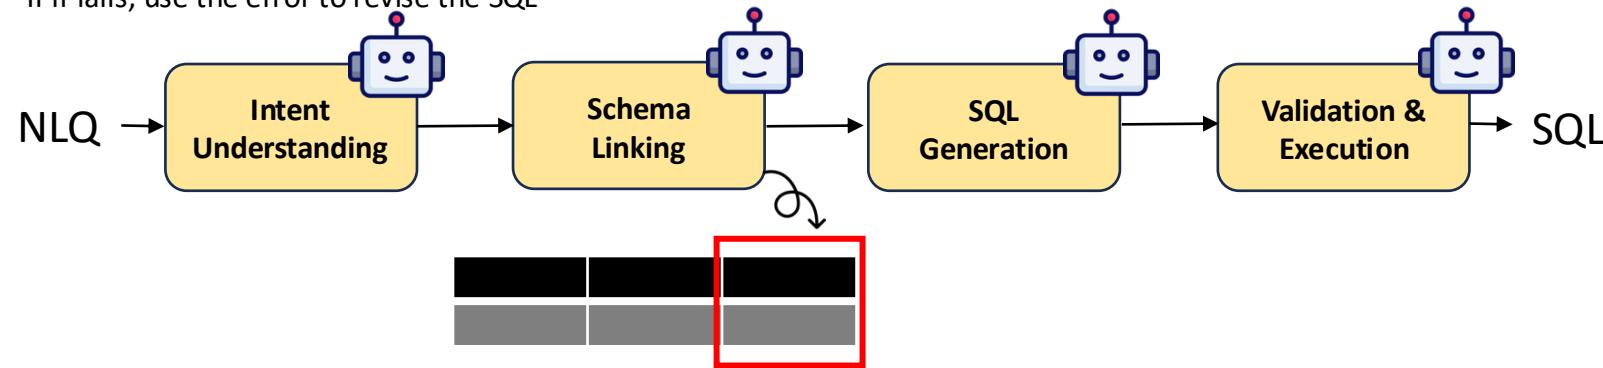
- Map words to tables/columns (and find the right join path)

## 3. SQL Generation

- Produce a valid SQL query (SELECT, JOIN, GROUP BY, HAVING, ...)

## 4. Validation & Execution

- Check syntax/types, run the query, and show results
- If it fails, use the error to revise the SQL



# Typical NL2SQL pipeline

**Intent Understanding:** What does the user really want? (filters, aggregation, time range, top-k, etc)

## Example

You are an analyst that extracts the user's intent from a natural-language database question.  
Do NOT write SQL.

Your job:

- 1) Restate the question in a clear and unambiguous way.
- 2) Identify the main task type (e.g., lookup, filter, aggregation, ranking, comparison, join).
- 3) Extract constraints (filters), metrics, grouping, sorting, and output columns.
- 4) If anything is ambiguous, list the ambiguity and propose 1–2 reasonable assumptions.
- 5) Produce the result in the JSON format below, with short strings.

Return ONLY valid JSON. No extra text.

Question: {NLQ}

## Typical NL2SQL pipeline

**Schema Linking:** Map words to tables/columns (and find the right join path)

### Example 1.

```
CREATE TABLE DEPT (DEPTNO NUMBER(2) CONSTRAINT PK_DEPT PRIMARY  
KEY, DNAME VARCHAR2(14) , LOC VARCHAR2(13) );
```

```
CREATE TABLE EMP (EMPNO NUMBER(4) CONSTRAINT PK_EMP PRIMARY KEY,  
ENAME VARCHAR2(10), JOB VARCHAR2(9), MGR NUMBER(4), HIREDATE DATE,  
SAL NUMBER(7,2), COMM NUMBER(7,2), DEPTNO NUMBER(2) CONSTRAINT  
FK_DEPTNO REFERENCES DEPT);
```

### Example 2.

```
DEPT (DEPTNO, DNAME, LOC)  
EMP (EMPNO, ENAME, JOB, MGR, HIREDATE, SAL, COMM, DEPTNO)
```

# Typical NL2SQL pipeline

**Validation & Execution:** Check syntax/types, run the query, and show results

- SQL parse/compile check, execution-based validation, retry, explain/query plan feedback, debugger style traces

## Example

You generated a SQL query for SQLite, but it failed.

Rules:

- Return exactly ONE SQLite SELECT query.
- Use only the given schema.
- Keep it minimal (avoid unnecessary joins).
- Add LIMIT 50 if no limit is present.

User question:

{NLQ}

Schema:

{SCHEMA}

Your SQL:

{SQL}

SQLite error:

{ERROR}

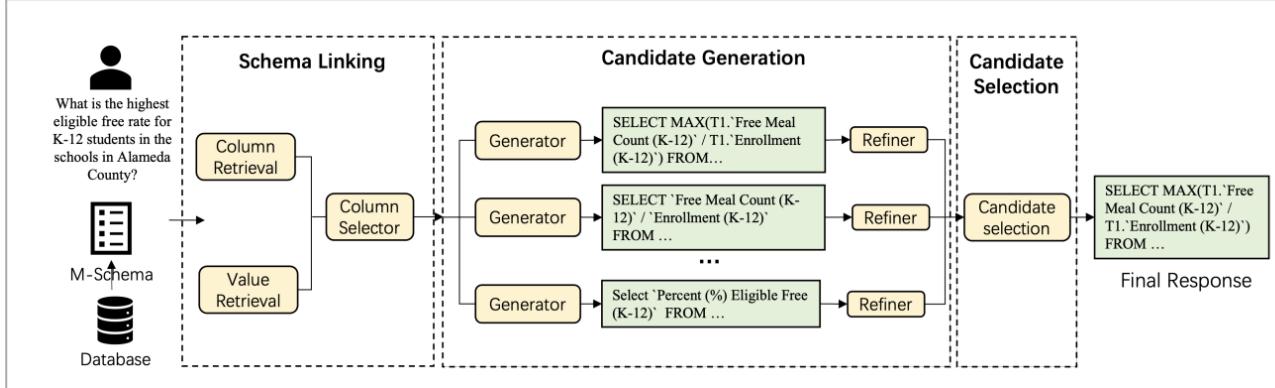
Output:

1) Fixed SQL only

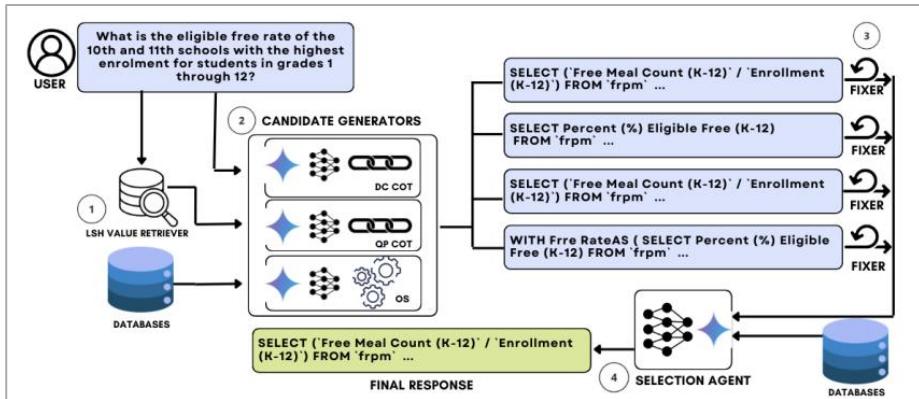
2) One-sentence reason for the fix

# Current NL2SQL Frameworks

- Academic research has developed pipelines that integrate diverse techniques.



XiYan-SQL: A Multi-Generator Ensemble Framework for Text-to-SQL



CHASE-SQL: Multi-Path Reasoning and Preference Optimized Candidate Selection in Text-to-SQL

## NL2SQL Platform

- Recently, many platforms that support NL2SQL have emerged.
  - <https://www.eversql.com/text-to-sql/>
  - <https://www.text2sql.ai>
  - <https://www.sqlai.ai/text-to-sql-ai>
- Major DBMS vendors (e.g., Oracle) are also developing NL2SQL capabilities.
  - <https://blogs.oracle.com/machinelearning/post/introducing-natural-language-to-sql-generation-on-autonomous-database>
- Snowflake introduces a semantic layer that abstracts schema objects into business-friendly concepts that are easier for humans to read and interpret.
  - <https://docs.snowflake.com/en/user-guide/views-semantic/overview>

## Hands-On

- In this lab, we will practice **NL2SQL (Natural Language → SQL)** using a Jupyter/Colab notebook.
  - We will use a **small sample database** (SCOTT-like tables) stored in **SQLite (in-memory)**.
  - You will write a few SQL queries manually, then try NL2SQL where an LLM generates SQL from your question.
- 
- **Prerequisite**
    - **Google Colab** (<https://colab.google/>)
    - No database installation needed: we use **SQLite in memory** inside the notebook
    - **OpenAI API key** : We will provide a key for this lab (do not share it publicly)